




Computer Graphics and Image Processing

Part 3: Image Processing
4 – Image Filtering

Martin Urschler, PhD


Intensity mapping vs. spatial filtering

- Up to now: Focus on greyvalue transformations
 - Modify contrast and brightness through linear transformation
 - Nonlinear transformation for histogram equalization
 - Input and output is a **single pixel**
- Spatial image filtering operates on a **neighborhood** (moving window)
 - Intensity transformations can be seen as operating on a 1×1 neighborhood
 - Image filtering: $n \times n$ neighborhood ($n > 1$, often n is odd!)
 - Again, linear or nonlinear transformations possible



Spatial image filtering

Example: Mean filter 5 x 5



Input:
235 x 235 px

Smoothing =
noise filtering




FIGURE 3.28 The mechanics of linear spatial filtering using a 3 × 3 filter mask. The form chosen to denote the coordinates of the filter mask coefficients simplifies writing expressions for linear filtering.

Basic principles of filtering

- Output filtered image $g(x, y)$ is obtained from input image $f(x, y)$ by applying **moving window transform** in $(2k + 1) \times (2l + 1)$ neighborhood (window): $g = MWT(f)$
 - Value $g(x, y)$ at each pixel location (x, y) is a certain linear or non-linear function of values of the original image $f(x, y)$
 - The values of f are taken in a $(2k + 1) \times (2l + 1)$ rectangle being centered on pixel location (x, y)
 - E.g. $(k = 1, l = 1)$ for a 3×3 window and $(k = 1, l = 3)$ for a 3×7 window
 - A general **linear MWT** multiplies each filter coefficient $w(\xi, \eta)$ with the image value that lies directly beneath it, and sums up these terms.
 - Example: Mean filter is just a sum with fixed weight!

$$\mu(x, y) = \frac{1}{(2k + 1)(2l + 1)} \sum_{\xi=-k}^k \sum_{\eta=-l}^l f(x + \xi, y + \eta)$$


Basic principles of filtering

- Windows are centered at each location (x, y)
- Popular windows – rectangles with odd sizes, e.g. 3×3 , 5×5 , 3×1 , etc.
- The 3×3 square window gives a set of offsets representing the neighborhood:

$$\{(\xi, \eta) : \xi = -1, 0, 1; \eta = -1, 0, 1\}$$

- The $(2k + 1) \times (2l + 1)$ rectangular window gives a set of offsets:

$$\{(\xi, \eta) : \xi = -k, \dots, -1, 0, 1, \dots, k; \eta = -l, \dots, -1, 0, 1, \dots, l\}$$



Spatial image filtering

- Example: Mean filter 5×5 applied to f


- Rectangular window ($k = 2, l = 2$):

$$\{(\xi, \eta) : \xi = -2, -1, 0, 1, 2; \eta = -2, -1, 0, 1, 2\}$$

- Mean filter generating output g :


$$g(x, y) = \mu(x, y) = \frac{1}{25} \sum_{\xi=-2}^2 \sum_{\eta=-2}^2 f(x + \xi, y + \eta)$$

$$w(\xi, \eta) = 1$$



$f(x, y)$

$g(x, y)$



Spatial image filtering

- Example: Mean filter 5×5 applied to f


- Rectangular window ($k = 2, l = 2$):

$$\{(\xi, \eta) : \xi = -2, -1, 0, 1, 2; \eta = -2, -1, 0, 1, 2\}$$

- Mean filter generating output g :


$$g(x, y) = \mu(x, y) = \frac{1}{25} \sum_{\xi=-2}^2 \sum_{\eta=-2}^2 f(x + \xi, y + \eta)$$

$$w(\xi, \eta) = 1$$



$f(x, y)$

$g(x, y)$



Spatial image filtering

- Example: Mean filter 5×5 applied to f


- Rectangular window ($k = 2, l = 2$):

$$\{(\xi, \eta) : \xi = -2, -1, 0, 1, 2; \eta = -2, -1, 0, 1, 2\}$$

- Mean filter generating output g :

$$g(x, y) = \mu(x, y) = \frac{1}{25} \sum_{\xi=-2}^2 \sum_{\eta=-2}^2 f(x + \xi, y + \eta)$$

$$w(\xi, \eta) = 1$$



$f(x, y)$

$g(x, y)$



Spatial image filtering

- Generic linear MWT:

$$g(x, y) = \frac{1}{(2k + 1)(2l + 1)} \sum_{\xi=-k}^k \sum_{\eta=-l}^l w(\xi, \eta) f(x + \xi, y + \eta)$$


- Mean filter 3 x 3:

$$w(\xi, \eta) = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- Vertical edge filter 3 x 3:

We will come back
to edges later!

$$w(\xi, \eta) = \frac{1}{4} \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$



Smoothing: Mean filter

■ Mean, box or average filtering:


- “Smoothing” images by reducing the variation of intensities between neighboring pixels -> **reduce white (Gaussian) noise**
- Each value is replaced with the average value of the neighboring pixels, including itself, normalized by window size to stay inside original greyvalue range!

■ Potential problems:

- A single “outlier” can significantly affect the average of all the pixels in a neighborhood
- Edge blurring: when the moving window crosses an edge, the filter will interpolate new pixel values on the edge
 - This may be a problem if sharp output edges are required!

Smoothing: Mean filter

■ One dimensional (1D) example



Moving window		Computing the mean		Rounding the mean and assigning to $g(x)$
$W_0 : (3, 3, 9)$	\Rightarrow	$(3 + 3 + 9)/3 = 5$	\Rightarrow	$g(0) = \text{round}\{5\} = 5$
$W_1 : (3, 9, 4)$	\Rightarrow	$(3 + 9 + 4)/3 = 5.33$	\Rightarrow	$g(1) = \text{round}\{5.33\} = 5$
$W_2 : (9, 4, 52)$	\Rightarrow	$(9 + 4 + 52)/3 = 21.67$	\Rightarrow	$g(2) = \text{round}\{21.67\} = 22$
$W_3 : (4, 52, 3)$	\Rightarrow	$(4 + 52 + 3)/3 = 19.67$	\Rightarrow	$g(3) = \text{round}\{19.67\} = 20$
$W_4 : (52, 3, 8)$	\Rightarrow	$(52 + 3 + 8)/3 = 21$	\Rightarrow	$g(4) = \text{round}\{21\} = 21$
$W_5 : (3, 8, 6)$	\Rightarrow	$(3 + 8 + 6)/3 = 5.67$	\Rightarrow	$g(5) = \text{round}\{5.67\} = 6$
$W_6 : (8, 6, 2)$	\Rightarrow	$(8 + 6 + 2)/3 = 5.33$	\Rightarrow	$g(6) = \text{round}\{5.33\} = 5$
$W_7 : (6, 2, 2)$	\Rightarrow	$(6 + 2 + 2)/3 = 3.33$	\Rightarrow	$g(7) = \text{round}\{3.33\} = 3$
$W_8 : (2, 2, 9)$	\Rightarrow	$(2 + 2 + 9)/3 = 4.33$	\Rightarrow	$g(8) = \text{round}\{4.33\} = 4$
$W_9 : (2, 9, 9)$	\Rightarrow	$(2 + 9 + 9)/3 = 6.67$	\Rightarrow	$g(9) = \text{round}\{6.67\} = 7$

g

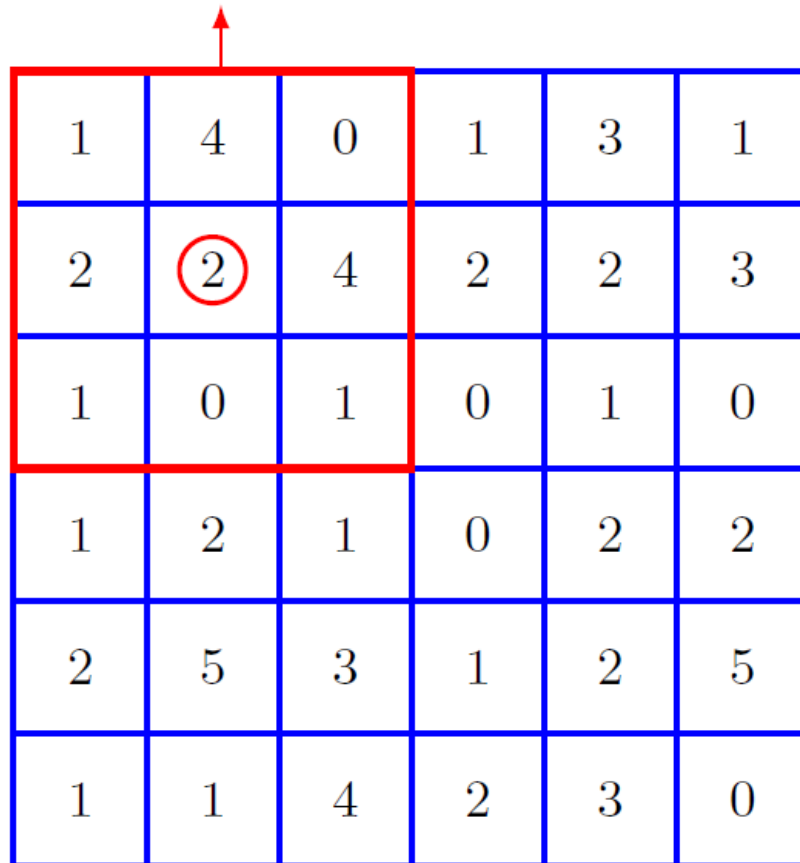
5	5	22	20	21	6	5	3	4	7
---	---	----	----	----	---	---	---	---	---

For $g(0)$ and $g(9)$, $f(0)$ and $f(9)$, respectively, are extended outside the boundaries.

2D Mean filter: 3 x 3 window

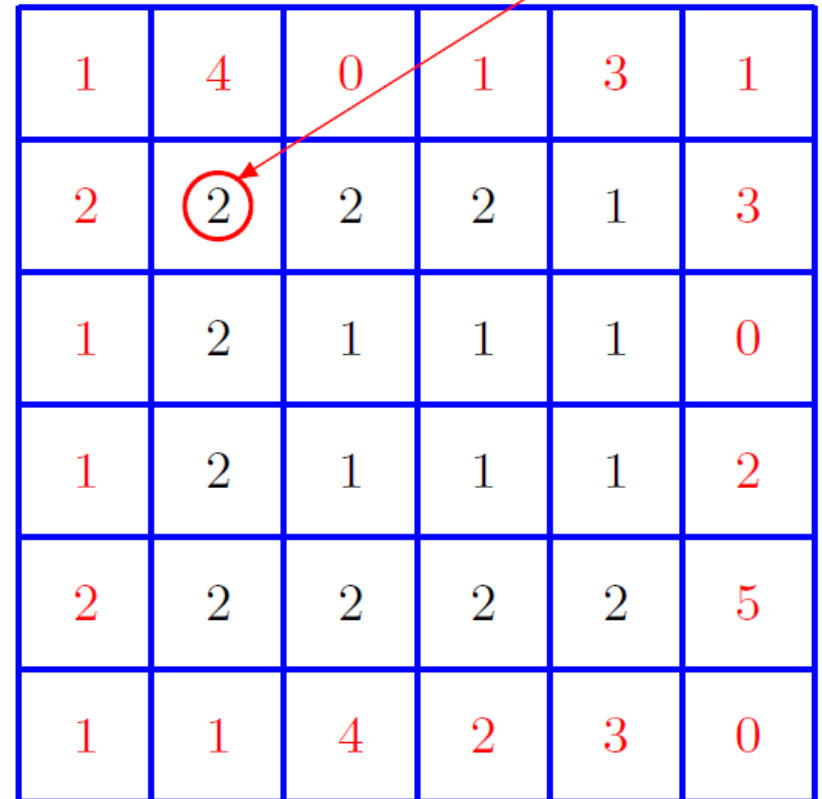
1: Keeping border values unchanged.

Averaging: $\text{round}\{(1 + 4 + 0 + 2 + 2 + 4 + 1 + 0 + 1)/9\} = 2$



1	4	0	1	3	1
2	2	4	2	2	3
1	0	1	0	1	0
1	2	1	0	2	2
2	5	3	1	2	5
1	1	4	2	3	0

Input f



1	4	0	1	3	1
2	2	2	2	1	3
1	2	1	1	1	0
1	2	1	1	1	2
2	2	2	2	2	5
1	1	4	2	3	0

Output g

2D Mean filter: 3 x 3 window

2: Extending **border values** outside with the boundary values.

$$\text{Averaging: } \text{round}\{(1 + 1 + 4 + 2 + 2 + 2 + 1 + 1 + 0)/9\} = 2$$

1	1	4	0	1	3	1	1
1	1	4	0	1	3	1	1
2	2	2	4	2	2	3	3
1	1	0	1	0	1	0	0
1	1	2	1	0	2	2	2
2	2	5	3	1	2	5	5
1	1	1	4	2	3	0	0
1	1	1	4	2	3	0	0

Input f

2	2	2	2	2	2
2	2	2	2	1	2
1	2	1	1	1	2
2	2	1	1	1	2
2	2	2	2	2	2
2	2	3	3	2	2

Output g

2D Mean filter: 3 x 3 window

3: Extending **border values** outside with zeros (zero padding).

Averaging: $\text{round}\{(0 + 1 + 4 + 0 + 2 + 2 + 0 + 1 + 0)/9 = 1$

0	0	0	0	0	0	0	0
0	1	4	0	1	3	1	0
0	2	2	4	2	2	3	0
0	1	0	1	0	1	0	0
0	1	2	1	0	2	2	0
0	2	5	3	1	2	5	0
0	1	1	4	2	3	0	0
0	0	0	0	0	0	0	0

Input f


1	1	1	1	1	1
1	2	2	2	1	1
1	2	1	1	1	1
1	2	1	1	1	1
1	2	2	2	2	2
1	2	2	2	1	1

Output g


2D Mean filter: more examples

FIGURE 3.33 (a) Original image, of size 500×500 pixels. (b)–(f) Results of smoothing with square averaging filter masks of sizes $m = 3, 5, 9, 15,$ and 35 , respectively. The black squares at the top are of sizes $3, 5, 9, 15, 25, 35, 45,$ and 55 pixels, respectively; their borders are 25 pixels apart. The letters at the bottom range in size from 10 to 24 points, in increments of 2 points; the large letter at the top is 60 points. The vertical bars are 5 pixels wide and 100 pixels high; their separation is 20 pixels. The diameter of the circles is 25 pixels, and their borders are 15 pixels apart; their intensity levels range from 0% to 100% black in increments of 20% . The background of the image is 10% black. The noisy rectangles are of size 50×120 pixels.

*Images from: Gonzalez & Woods,
Digital Image Processing, 3rd ed.*



Noisy X-ray image
of circuit board




Noise reduction with
 3×3 averaging filter



2D Mean filter: Separability


- For larger filter windows (think e.g. 15 x 15), two nested for loops are increasingly inefficient.
- **Separable filter windows** can be implemented more efficiently by composition of two operations with 1D filters!



$$w(\xi, \eta) = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



$O(k^2)$ operations
per pixel



$$w_{row}(\xi, \eta) = \frac{1}{3} [1 \ 1 \ 1]$$




$$w_{col}(\xi, \eta) = \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$



$O(k)$ operations per pixel, but needs memory for intermediate image!

Gaussian linear filter

- To blur images and remove noise and fine detail
- One-dimensional Gaussian function (zero mean)
- The pixels further away contribute less weight...




$$G(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}}$$

σ – standard deviation

$\kappa = \frac{x}{\sigma}$	0	1	2
$G(\kappa)$	0.399	0.242	0.054
$\frac{G(\kappa)}{G(0)}$	1.000	0.607	0.135

Gaussian linear filter

- Standard deviation σ of Gaussian probability density function guides its behaviour:
 - 68% of the x-values are in the range $[mean - \sigma, mean + \sigma]$
 - 95% of the x-values are in the range $[mean - 2\sigma, mean + 2\sigma]$
 - 99.7% of the x-values are between $[mean - 3\sigma, mean + 3\sigma]$



1D Gaussian linear filter

- Built using normalized 1D Gaussian $G(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}}$
- Filter kernel $s(x)$: a discrete approximation of the normalized Gaussian; e.g. for the 1×5 filter with $\sigma = 1$:

x	-2	-1	0	1	2
$s(x)$	0.054	0.242	0.399	0.242	0.054
Old: $[10s(x)]$	1	2	4	2	1
Old: kernel	$\frac{1}{10}$	$\frac{2}{10}$	$\frac{4}{10}$	$\frac{2}{10}$	$\frac{1}{10}$

- Rule of thumb: set filter half-width to about $2-3\sigma$
- Still often used filter kernel for 1×3 filter with $\sigma = 1$ is $\frac{1}{4} [1 \ 2 \ 1]$


Gaussian linear filter

- 2D Gaussian filter is built using the (isotropic) 2D Gaussian function

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

- 2D Gaussian is separable!

$$G(x, y) = G(x)G(y)$$




2D Gaussian linear filter


- Filter kernel $s(x)$: a discrete approximation of the normalized continuous 2D Gaussian; e.g. for the 3×3 and 5×5 filters with $\sigma = 1$:

$$s_{3 \times 3}(x) = \frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \quad s_{5 \times 5}(x) = \frac{1}{100} \begin{bmatrix} 1 & 2 & 4 & 2 & 1 \\ 2 & 4 & 8 & 4 & 2 \\ 4 & 8 & 16 & 8 & 4 \\ 2 & 4 & 8 & 4 & 2 \\ 1 & 2 & 4 & 2 & 1 \end{bmatrix}$$

- The larger the value of σ , the wider the peak of the Gaussian and the larger the blurring
- Non-uniform averaging: low pass filtering
- Rotational symmetry with no directional bias
- Fast computations due to separability
- Might not preserve image brightness





$\sigma = 2$; 30×30 kernel




$\sigma = 5$; 30×30 kernel

Gaussian vs. mean (box) filter



Gaussian filter for scale space

Orig



3x3




5x5




11x11



Gaussian filter with Salt-and-pepper noise




Noisy input image



Gaussian filter: blurred edges, residual noise

Median filter



Noisy input image



Median filter output (5x5)

Median filter


- Effective **nonlinear** filter, being used frequently to remove “salt-and-pepper” noise while preserving edges
- Replace each pixel with the median of the neighborhood of K pixels:

$$g(x, y) = \text{median}\{f(x + \xi_i, y + \eta_i) : i = 1, \dots, K\}$$


- Computation for each pixel (x, y) :
 1. Sort all K values from the neighboring window of $f(x, y)$ into ascending numerical order: $v_{[0]} \leq v_{[1]} \leq \dots \leq v_{[K-1]}$
 2. Select the middle value, $g(x, y) = v_{[K/2]}$, if K is odd or the average, $g(x, y) = 0.5(v_{[(K-1)/2]} + v_{[(K+1)/2]})$, of the two middle values if K is even

Median filter: 1D example

Median filtering of a simple 1D signal $\mathbf{f} = (f(x) : x = 0, 1, \dots, 9)$ with a moving window $W_x = [x - 1, x, x + 1]$ of size 3.



Median filter: 1D example



Moving window		Sorting		Selecting the median and assigning to $g(x)$
W_0 : (3, 3, 9)	\Rightarrow	(3, 3 , 9)	\Rightarrow	$g(0) = \text{median}\{3, 3, 9\} = 3$
W_1 : (3, 9, 4)	\Rightarrow	(3, 4, 9)	\Rightarrow	$g(1) = \text{median}\{3, 9, 4\} = 4$
W_2 : (9, 4, 52)	\Rightarrow	(4, 9 , 52)	\Rightarrow	$g(2) = \text{median}\{9, 4, 52\} = 9$
W_3 : (4, 52, 3)	\Rightarrow	(3, 4, 52)	\Rightarrow	$g(3) = \text{median}\{4, 52, 3\} = 4$
W_4 : (52, 3, 8)	\Rightarrow	(3, 8 , 52)	\Rightarrow	$g(4) = \text{median}\{52, 3, 8\} = 8$
W_5 : (3, 8, 6)	\Rightarrow	(3, 6 , 8)	\Rightarrow	$g(5) = \text{median}\{3, 8, 6\} = 6$
W_6 : (8, 6, 2)	\Rightarrow	(2, 6 , 8)	\Rightarrow	$g(6) = \text{median}\{8, 6, 2\} = 6$
W_7 : (6, 2, 2)	\Rightarrow	(2, 2 , 6)	\Rightarrow	$g(7) = \text{median}\{6, 2, 2\} = 2$
W_8 : (2, 2, 9)	\Rightarrow	(2, 2 , 9)	\Rightarrow	$g(8) = \text{median}\{2, 2, 9\} = 2$
W_9 : (2, 9, 9)	\Rightarrow	(2, 9 , 9)	\Rightarrow	$g(9) = \text{median}\{2, 9, 9\} = 9$



For $g(0)$ and $g(9)$, $f(0)$ and $f(9)$, respectively, are extended outside the boundaries.

Median filter

Keeping **border values** unchanged (processing no border pixels).

Sorted: (0, 0, 1, 1, **1**, 2, 2, 4, 4)

1	4	0	1	3	1
2	2	4	2	2	3
1	0	1	0	1	0
1	2	1	0	2	2
2	5	3	1	2	5
1	1	4	2	3	0


Input f

1	4	0	1	3	1
2	1	1	1	1	3
1	1	1	1	2	0
1	1	1	1	1	2
2	2	2	2	2	5
1	1	4	2	3	0

Output g


Median filter

1: Extending **border values** outside with the boundary values.



Median filter

2: Extending border values outside with zeros.



Gauss vs. median filter

3×3 GF



5×5 GF



11×11 GF



3×3 MF




5×5 MF



11×11 MF

Gauss vs. median filter



Gauss vs. median filter

- Use depends on noise assumption
 - White noise: mean or Gaussian filter
 - Salt-and-pepper noise: median filter (robust to outliers!)
- Local intensity transitions
 - Remain unchanged with mean or Gaussian filter
 - Can get destroyed by median filter
 - Gauss filter: Important for Gaussian scale space (multi-scale representation to inspect image at different complexity levels)