

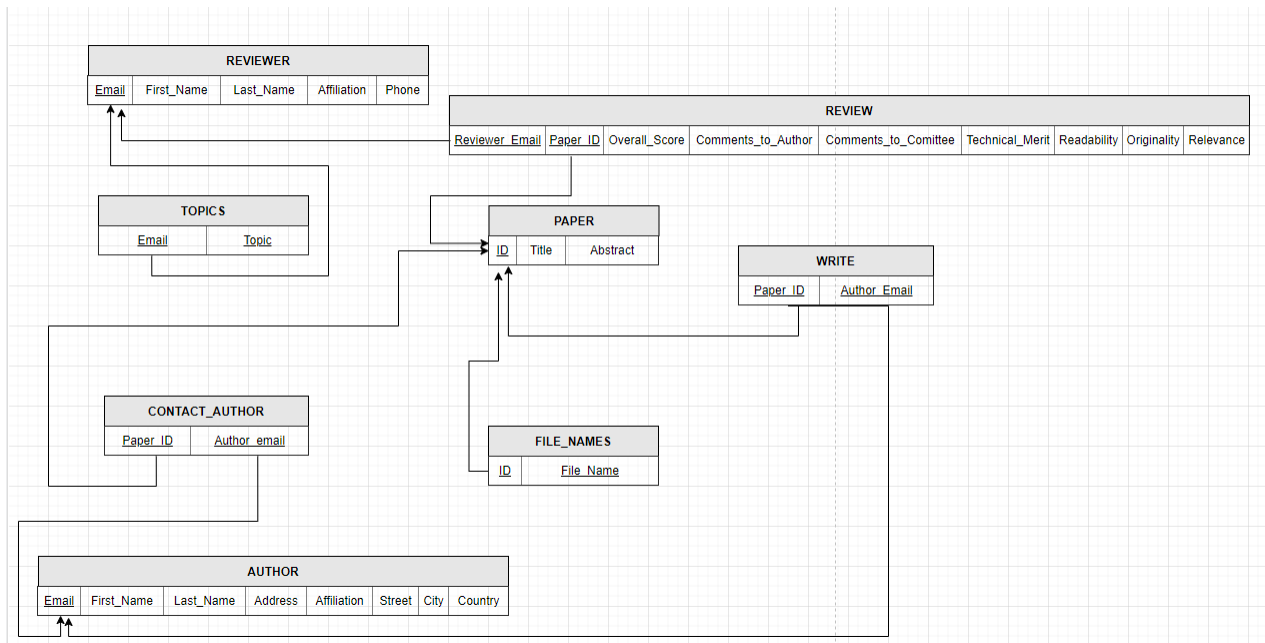
SOFTENG 351: Exam

July 4, 2020

Aiden Burgess

abur970 - 600280511

Question 1



Question 2

a)

i.

```
ALTER TABLE COURSE ALTER Department SET DEFAULT 'CS';
```

ii.

```
ALTER TABLE STUDENT ADD COLUMN Birth_Date DATE;
```

iii.

```
UPDATE GRADE_REPORT
SET Grade='B'
WHERE Student_number='17' AND Section_identifier='119';
```

iv.

```
SELECT Course_name, Course_number, Credit_hours, Semester, Year, Grade
FROM (GRADE_REPORT NATURAL JOIN SECTION NATURAL JOIN COURSE)
WHERE Student_number='8'
```

b)

i.

$$ANDERSON_SECTIONS \leftarrow \sigma_{Instructor='Anderson'}(SECTION)$$

$$NUM_STUDENTS \leftarrow Section_identifier \Join_{COUNT Student_number} (GRADE_REPORT * ANDERSON_SECTIONS)$$

$$COURSE_INFO \leftarrow \pi_{Course_name, Course_number, Semester, Year}(COURSE * ANDERSON_SECTIONS)$$

$$RESULT \leftarrow COURSE_INFO * NUM_STUDENTS$$

ii.

Assume a student can not take two of the same course in the same year.

Assume that to satisfy the requirement a student needs to take all courses (non duplicate) available in the year '08', regardless of major.

$$ALL_COURSES \leftarrow \pi_{Course_number}(\sigma_{Year='08'}(SECTION))$$

$$STUDENT_COURSES \leftarrow \sigma_{Year='08'}(STUDENT * GRADE_REPORT * SECTION)$$


$$NUM_COURSES(Total_courses) \leftarrow \Join_{COUNT Course_number} (ALL_COURSES)$$


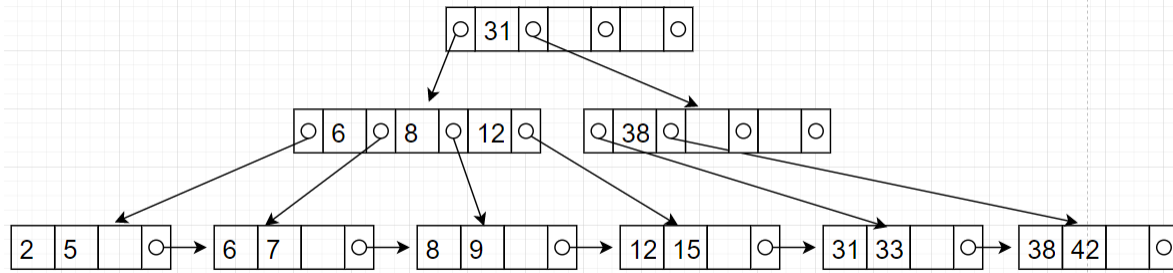
$$STUDENT_NUM_COURSES(Student_courses) \leftarrow Student_number \Join_{COUNT Course_number} (STUDENT_COURSES)$$

$$RESULT(Name, Major) \leftarrow NUM_COURSES \Join_{Total_courses=Student_courses} STUDENT_NUM_COURSES$$

Question 3

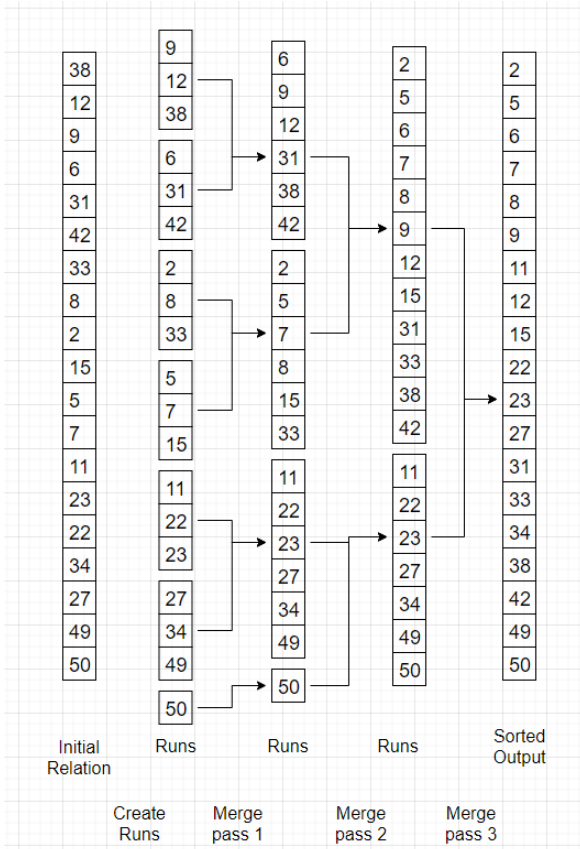
a)





b)


i.



Each pass takes $19 \times 2 = 38 I/Os$

As there are three passes, the sorting takes $38 \times 3 = 114 I/Os$ in total.

ii.



Question 4

Assume that each attribute of a “combined key” is a key as well.

Nested Loop Join then Nested Loop Join

Assume that joining the smallest relations first leads to lowest cost. Therefore, join y and z.

As z is smaller than y, it will be used in the outer loop.

$$y \bowtie z = |z| \times \frac{n_y}{\max(V(C, y), V(C, z))} = 300 \times \frac{200}{100} = 300 I/Os$$

As y and z have a common key of A, the calculation of $n_{y \bowtie z}$ is relatively simple:

$$|y \bowtie z| = \min(|y|, |z|) = 300$$

As there is no pipelining, the intermediate result must be written:

$$\text{Cost of writing} = \frac{|y \bowtie z|}{bfr} = \frac{300}{2} = 150 I/Os$$

As $y \bowtie z$ is smaller than x, it will be used in the outer loop.

$$x \bowtie (y \bowtie z) = |y \bowtie z| \times \frac{n_x}{V(A, x)} = 300 \times \frac{250}{50} = 1500 I/Os$$

$$\text{Total Cost} = 300 + 150 + 1500 = 1950 I/Os$$

Merge Join then Nested Loop Join

As x and y are both clustered on B, then

$$x \bowtie y = n_x + n_y = 250 + 200 = 450 \text{ I/Os}$$

As x and y have a common key of B, the calculation of $n_{x \bowtie y}$ is relatively simple:

$$|x \bowtie y| = \min(|x|, |y|) = 200$$

As there is no pipelining, the intermediate result must be written:

$$\text{Cost of writing} = \frac{|x \bowtie y|}{bfr} = \frac{200}{2} = 100 \text{ I/Os}$$

$$(x \bowtie y) \bowtie z = n_{x \bowtie y} + n_z = 200 + 400 = 600 \text{ I/Os}$$

$$\text{Total Cost} = 450 + 100 + 600 = 1150 \text{ I/Os}$$

Therefore, the most optimal cost is by using merge join on x and z, then using nested loop join on the resultant relation with y.


The cost is 1150 I/Os.

Question 5

a)

Deadlock occurs when a cycle appears in a wait-for graph. This means that a transaction, T_1 is waiting for a lock held by another transaction T_2 , and that neither can be completed until the other has been completed. However, this is not possible using conservative 2PL because the earlier transaction T_1 which started earlier has already requested the locks needed, so it does not wait for a lock held by T_2 . By induction, this process can be repeated until we reach the oldest transaction in the schedule. This transaction T_0 does not wait for any locks and therefore completes successfully and releases all locks held. This process is repeated until all transactions are completed. Therefore, a cycle is not possible, as there is an order to which locks are requested such that the oldest transaction does not wait for a lock held by any other.


b)



As there is a cycle in the graph (e.g. T_1 and T_2), the schedule is not conflict-serializable.

c)

i.



There is a deadlock as a cycle exists between T_1 and T_2 .

ii.

T_1 is waiting on T_2 for the lock on B.

T_2 is waiting on T_1 for the lock on A.

As T_2 has a higher TS, T_2 is aborted, and subsequently the lock on A is released. Immediately, T_1 grabs the lock exclusive lock on B and T_2 is waiting for T_1 for the lock on A and B. This is accepted because $TS(T_1) < TS(T_2)$.

This allows T_1 to finish its transaction and commit. Subsequently, T_2 can grab the locks for A and B and complete its transaction.

Therefore, the deadlock has been avoided.

Question 6

1

As the transaction has been committed and is included in a checkpoint, no further instructions are needed.
No action

2

Redo all write_item operations of T_2 from the log after the checkpoint, in the order in which they were written into the log, using the REDO procedure.

3

Undo all the write_item operations of T_3 using the UNDO procedure before the checkpoint. The operations should be undone in the reverse of the order in which they were written to the log.

4

Redo all write_item operations of T_4 from the log after the checkpoint, in the order in which they were written into the log, using the REDO procedure.

5

Undo all the write_item operations of T_5 using the UNDO procedure after the checkpoint. The operations should be undone in the reverse of the order in which they were written to the log.