

# **SOFTENG 351: Lab #10**

June 4, 2020

**Aiden Burgess**

abur970 - 600280511

## 1.

No indices, so we can rule out hash join immediately. Can not use merge join as it requires ordering on the attributes.

Therefore, we must use block nested-loop join which will require:

$$n_s + n_r \text{ IOs}$$

Where  $n_r$  is the number of blocks  $r$  is stored in,  $n_s$  is the number of blocks  $s$  is stored in.

The amount of memory needed is at least  $\min(n_s, n_r) + 1$  blocks in main memory as we need 1 block free for the output buffer.

## 2.

Simple translation:  $\pi_{T.branch\_name}(\sigma_{T.assets > S.assets \wedge S.branch\_city = "Brooklyn"}((\rho_T(branch) \times (\rho_S(branch))))$

Optimized:  $\pi_{T.branch\_name}(\sigma_{T.assets > S.assets}(\rho_T(\pi_{branch\_name, assets}(branch)) \times \rho_S(\pi_{assets}(\sigma_{branch\_city = "Brooklyn"}))))$

In the optimized version, the selection and projection statements have been moved down as far as possible, so that operations higher in the hierarchy do not perform as much work.

## 3.

a)

## SQL result

**Host:** Student mysql database

**Database:** stu\_abur970\_SOFTENG\_351\_C\_S1\_2020\_A2\_Q1

**Generation Time:** Jun 03, 2020 at 08:43 PM

**Generated by:** phpMyAdmin 4.0.10.7 / MySQL 5.1.73-log

**SQL query:** EXPLAIN SELECT \* FROM DEPARTMENT WHERE Dname="Headquarters";

**Rows:** 1

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	DEPARTMENT	ALL	NULL	NULL	NULL	NULL	6	Using where

b)

## SQL result

**Host:** Student mysql database

**Database:** stu\_abur970\_SOFTENG\_351\_C\_S1\_2020\_A2\_Q1

**Generation Time:** Jun 03, 2020 at 08:30 PM

**Generated by:** phpMyAdmin 4.0.10.7 / MySQL 5.1.73-log

**SQL query:** EXPLAIN SELECT \* FROM DEPARTMENT WHERE Dname="Headquarters";

**Rows:** 1

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	DEPARTMENT	const	Dname	Dname	17	const	1	

c)

## SQL result

**Host:** Student mysql database

**Database:** stu\_abur970\_SOFTENG\_351\_C\_S1\_2020\_A2\_Q1

**Generation Time:** Jun 03, 2020 at 09:13 PM

**Generated by:** phpMyAdmin 4.0.10.7 / MySQL 5.1.73-log

**SQL query:** EXPLAIN SELECT \* FROM DEPARTMENT INNER JOIN EMPLOYEE ON Dnumber=Dno;

**Rows:** 2

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	DEPARTMENT	ALL	PRIMARY	NULL	NULL	NULL	6	
1	SIMPLE	EMPLOYEE	ref	Dno	Dno	4	stu_abur970_SOFTENG_351_C_S1_2020_A2_Q1.DEPARTMENT...	3	

d)

## SQL result

**Host:** Student mysql database

**Database:** stu\_abur970\_SOFTENG\_351\_C\_S1\_2020\_A2\_Q1

**Generation Time:** Jun 03, 2020 at 09:16 PM

**Generated by:** phpMyAdmin 4.0.10.7 / MySQL 5.1.73-log

**SQL query:** EXPLAIN SELECT \* FROM DEPENDENT INNER JOIN WORKS\_ON WHERE Hours>8;

**Rows:** 2

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	DEPENDENT	ALL	NULL	NULL	NULL	NULL	11	
1	SIMPLE	WORKS_ON	ALL	NULL	NULL	NULL	NULL	48	Using where; Using join buffer

**4.**

The size of  $r_1 \bowtie r_2 \bowtie r_3$  is estimated to be 1000 tuples, as C is a foreign key referencing  $r_2$ , and E is a foreign key referencing  $r_3$ .

The strategy for computing the join is to go through each tuple in  $r_1$  sequentially.

For each tuple (A, B, C):

- Lookup tuple for C in  $r_2$ . As C is the primary key of  $r_2$ , there can be at most one tuple. This can be done efficiently using an index.
- Lookup tuple for E in  $r_3$ . As E is the primary key of  $r_3$ , there can be at most one tuple. This can be done efficiently using an index.

**5.**

Consider the relations  $r_1(A, B, C)$ ,  $r_2(C, D, E)$ , and  $r_3(E, F)$  without primary keys, except the entire schema. Let  $V(C, r_1)$  be 900,  $V(C, r_2)$  be 1100,  $V(E, r_2)$  be 50, and  $V(E, r_3)$  be 100. Assume that  $r_1$  has 1000 tuples,  $r_2$  has 1500 tuples, and  $r_3$  has 750 tuples. Estimate the size of  $r_1 \bowtie r_2 \bowtie r_3$  and give an efficient strategy for computing the join.

$$\text{estimate}(|r_1 \bowtie r_2|) = \min(n_{r_2} \times \frac{n_{r_1}}{V(C, r_1)}, n_{r_1} \times \frac{n_{r_2}}{V(C, r_2)}) = \min(1667, 1364) = 1364 \text{ tuples}$$

$$\text{estimate}(|r_1 \bowtie r_2 \bowtie r_3|) = n_{r_1 \bowtie r_2} \times \frac{n_{r_3}}{V(E, r_3)} = 1364 \times \frac{750}{100} = 10230 \text{ tuples}$$

**6.****Atomicity**

Transactions are either performed entirely or not at all. This prevents errors where transactions are stopped in the middle. These errors can violate integrity constraints on the database and be difficult to rollback.

**Consistency Preservation**

A transaction should take a database from one consistent state to another. Again, this prevents the database from reaching an invalid state which violates constraints. For applications like banking, if an invalid state occurs, many users could be impacted.

**Isolation**

A transaction should not be interfered by other transactions. If another transaction interfered with the current transaction, calculations could be wrong, leading to invalid database states. This maintains the integrity of the data and helps with concurrency control.

**Durability/Permanency**

A transaction's changes must be persistent in the database. Once a transaction has been completed, it can not be completely removed from the history of the database. It prevents physical issues like power outages or hardware failure to affect the database state.