

# Query Processing and Optimization



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE

1. Introduction
2. Physical-Query-Plan Operators
3. Query Optimization

Reference:

Chapter 13-14, Database System Concepts, Sixth Edition  
Chapter 15, Database System, the Complete Book.

# Introduction

```
table course(course_id, title, dept_name);  
table instructor(ID, name, dept_name, salary);  
table teaches(ID, course_id, sec_id, semester, year);  
  
SELECT name, title  
FROM instructor  
      NATURAL JOIN teaches  
      NATURAL JOIN course  
WHERE dept_name='Music';
```

- Use the knowledge we have learnt this week to discuss:
  - What are the possible ways to answer the above SQL query?
  - Which one is better?

# Introduction

```
SELECT name, title
FROM instructor
    NATURAL JOIN teaches
    NATURAL JOIN course
WHERE dept_name='Music';
```



## Relational Algebra

- $\Pi_{name, title}(\sigma_{dept\_name=Music}(instructor \bowtie (teaches \bowtie course)))$
- $\Pi_{name, title}(\sigma_{dept\_name=Music}(instructor) \bowtie (teaches \bowtie course))$

# Introduction

- $\Pi_{name, title}(\sigma_{dept\_name=Music}(instructor \bowtie (teaches \bowtie course)))$
- $\Pi_{name, title}(\sigma_{dept\_name=Music}(instructor) \bowtie (teaches \bowtie course))$

## Query plan (Logic plan):

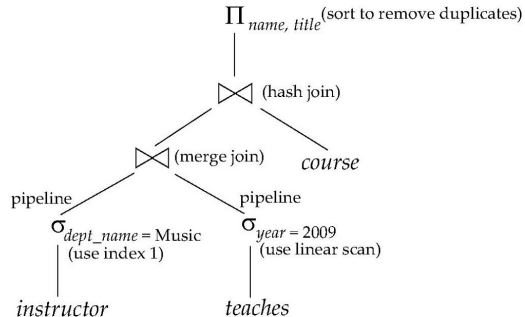


# Introduction

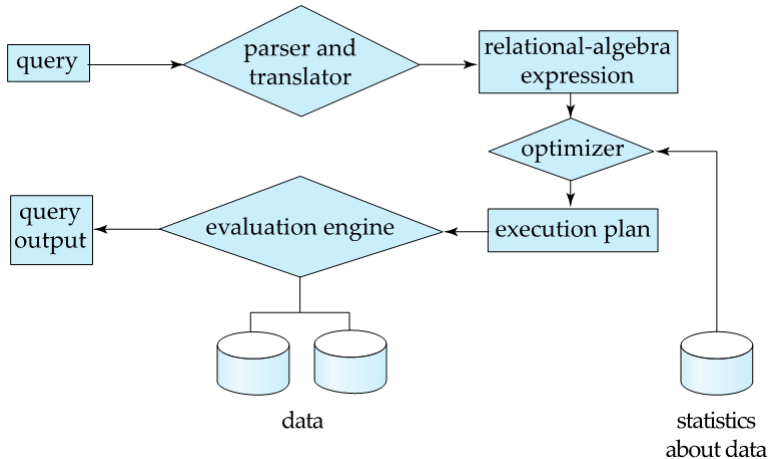
■  $\Pi_{name, title}((\sigma_{dept\_name=Music} instructor \bowtie \sigma_{year=2009} teaches) \bowtie course)$

## Physical/Evaluation plan

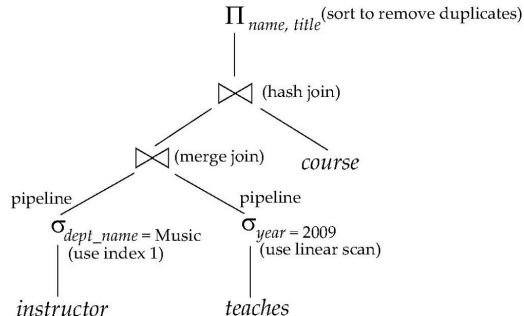
: an annotated expression on detailed evaluation strategy.



# Introduction



# Introduction



**Cost estimation** The cost of an operator is dependent on

- Metadata: The size of the underlying relation.
- Implementation: E.g., table scan, or using indexes (B+tree/hashing/...).
- reporting: Whether the result needs to be stored on disk.

7

In this course, we assume that we don't need to store the query results on the disk.

# Overview



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE

1. Introduction
2. Physical-Query-Plan Operators
  - One-pass algorithms
  - Multi-pass algorithms
  - Index-based algorithms



# Physical-Query-Plan Operators

## External Memory Model

Denote by  $B$  the block size of the system. Let  $M$  be the size of the main memory. The cost of an algorithm is decided by the number of blocks (# of I/Os) transferred between the main memory and the disk.

Sometimes we use  $m = M/B$  pages as the size of the main memory to simplify the analysis. In this scenario, the allocation of the memory is in pages.

# Physical-Query-Plan Operators

Iterators:

```
Open() {
    b := the first block of R;
    t := the first tuple of block b;
}

GetNext() {
    IF (t is past the last tuple on block b) {
        increment b to the next block;
        IF (there is no next block)
            RETURN NotFound;
        ELSE /* b is a new block */
            t := first tuple on block b;
    } /* now we are ready to return t and increment */
    oldt := t;
    increment t to the next tuple of b;
    RETURN oldt;
}

Close() {
}
```

# Physical-Query-Plan Operators



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE

Iterators:

- Combines several operations into one
- Avoids writing temporary files
- Many iterators may be active at one time

# Overview



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE


1. Introduction
2. Physical-Query-Plan Operators
  - One-pass algorithms
  - Multi-pass algorithms
  - Index-based algorithms

# One-pass algorithms

- Tuple-at-a-time, unary operations: selection and projection
- Full-relation, unary operations: the grouping operator and duplication deduction operator, conditioned that the relation can fit in to the main memory.
- Full-relation, binary operations: unions, intersection, difference, join, product, conditioned that one of the relation fit into the main memory.


# One-pass algorithms

- Tuple-at-a-time, unary operations: selection and projection



# One-pass algorithms

- Tuple-at-a-time, unary operations: selection and projection
- Full-relation, unary operations: the grouping operator and duplication deduction operator, conditioned that the relation can fit in to the main memory.



# One-pass algorithms

- Tuple-at-a-time, unary operations: selection and projection
- Full-relation, unary operations: the grouping operator and duplication deduction operator, conditioned that the relation can fit in to the main memory.
- Full-relation, binary operations: unions, intersection, difference, join, product, conditioned that one of the relation fit into the main memory.

Arrangement of the memory:

- $M - B - 1$ : Hold one relation
- $B$ : buffer one page of the other relation
- 1: buffer for the output



# Overview



1. Introduction
2. Physical-Query-Plan Operators
  - One-pass algorithms
  - Multi-pass algorithms
    - Multi-way merge sort
    - Join with both relations larger than the memory size
  - Index-based algorithms

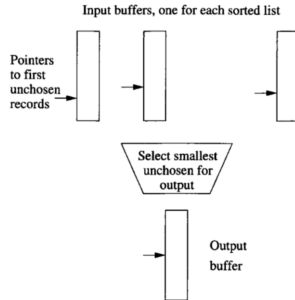
# Multi-way Merge Sort

Sort-based algorithms:

- Duplication elimination
- Grouping and aggregation
- Union
- Intersection and difference

# Multi-way Merge Sort

Atom operation: merge  $m - 1$  sorted list into one sorted list by allocating  $m - 1$  input buffer and 1 output buffer. The cost is linear to the size of the data.



## Multi-way Merge Sort

Sort the data of  $n$  blocks on a machine with a memory of  $m$  blocks.

- Create runs (pass 0):
    - Each time load  $m$  blocks into memory, sort them and then output to create a run
  - Merge pass  $i$ ,  $i = 1, 2, \dots$ : for every  $m - 1$  runs created by pass  $i - 1$ 
    - synchronize scan these runs to create a single run
    - one block used for output buffer
- until a pass creates only one run.

### Complexity.

$$n \log_{m-1} \frac{n}{m} + n \approx n \log_m n \text{ I/Os}$$

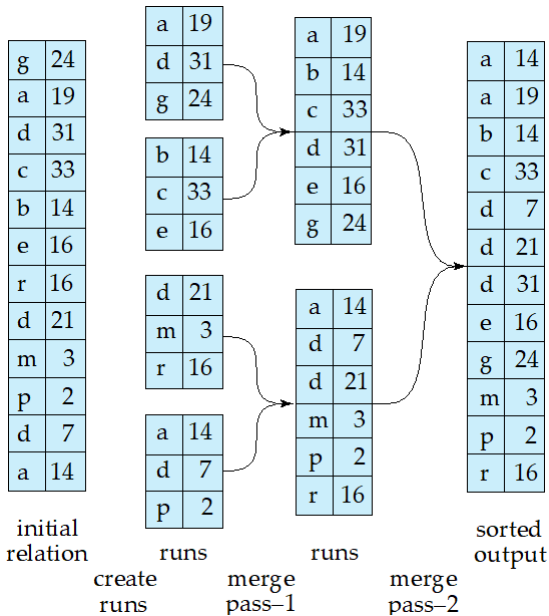


Figure: Example:  $m = 3$ , blocking factor = 1

# Join



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE

- Block nested-loop join
- Merge join
- Hash join

## Block nested-loop join

$r \bowtie s$ :  $r$  is stored in  $n_r$  blocks.  $s$  is stored in  $n_s$  blocks

```
for each chunk of  $m-1$  blocks of  $r$ 
  read these blocks into the main memory buffers;
  organize their tuples into a search structure whose search key is the
     $\hookrightarrow$  common attributes of  $r$  and  $s$ ;
  for each block  $bs$  of  $s$ 
    load  $bs$  into the main memory
    for each tuple  $t$  of  $bs$ 
      find the tuples of  $r$  in the buffer that join with  $t$ 
      report the result
```

**Complexity:**  $\lceil n_r / (m - 1) \rceil n_s + n_r$  I/Os.

# Merge Join

$r \bowtie s$ :  $r$  is stored in  $n_r$  blocks.  $s$  is stored in  $n_s$  blocks.

Join attribute:  $A$ .

Assume that  $r$  and  $s$  are both ordered  $A$ .

**for** each value  $x$  of attribute  $A$  shared by  $s$  **and**  $r$  (synchronize scan)  
    block-nested-loop join on subrelations  $s[A=x]$  **and**  $r[A=x]$

	$a1$	$a2$
$pr \rightarrow$	a	3
	b	1
	d	8
	d	13
	f	7
	m	5
	q	6
	$r$	

	$a1$	$a3$
$ps \rightarrow$	a	A
	b	G
	c	L
	d	N
	m	B
	$s$	

## Complexity.

- Best case:  $n_r + n_s$
- Worst case: block-nested-loop-join
- Consider the cost when
  - $A$  is a key of  $r$ , or
  - all the attribute values are having the same frequency

## Hashing: Partition a relation $r$ to $k$ buckets.

Consider a relation  $r$  with an attribute  $A$  and an integer  $k \in [1, m)$ ;  
Hash function  $h_A()$  maps a tuple  $t$  to an integer in range  $[0, k)$  based on its value on  $A$ .  
initialize  $k$  buckets using  $k$  empty buffers;  
**for** each block  $b$  of relation  $r$   
    read block  $b$  into the  $(k+1)$ -th buffer;  
    **for** each tuple  $t$  **in**  $b$   
        **if** the buffer **for** bucket  $h(t)$  has no room **for**  $t$   
            write the buffer to disk;  
            initialize a new empty block **in** that buffer;  
            copy  $t$  to the buffer **for** bucket  $h(t)$ ;  
**for** each bucket **do**  
    **if** buffer **for** this bucket **is not** empty  
        write the buffer to disk;

24 Denote by  $r_i$  the  $i$ -th bucket of  $r$ .




# Hash Join

$r \bowtie s$ :  $r$  is stored in  $n_r$  blocks.  $s$  is stored in  $n_s$  blocks.

Join attribute:  $A$ .

Hash function  $h_A()$  maps from a tuple to an integer range  $[0, k)$  based on  $A$ .

1. Partition  $r$  to  $k$  buckets using hash function  $h_A$
2. Partition  $s$  to  $k$  buckets using hash function  $h_A$
3. For each integer  $i \in [0, k)$ 
  - Join on  $r_i$  and  $s_i$  (one pass/multi pass)



## Examples: Cost Estimation

Goal is to count disk I/Os. But we First have to estimate sizes of intermediate results?

- $V(A,r)$  number of distinct values of attribute  $A$  in relation  $r$ .
- $|r|$  number of tuples in relation  $r$ .

Consider relation  $r(A, B)$  with  $n_r = 1000$  blocks and relation  $s(A, C)$  with  $n_s = 500$  blocks. The memory had  $m = 101$  pages.

- Nested-loop-join
- Multi-way merge join
- Hash join

## Examples: Cost Estimation, Nested-Loop-Join

Consider relation  $r(A, B)$  with  $n_r = 1000$  blocks and relation  $s(A, C)$  with  $n_s = 500$  blocks. The memory had  $m = 101$  pages.

- Outer-loop: load chunk of 100 blocks of  $s$  to the main memory  
5 chunks, 100 blocks each
- Inner-loop: scan  $r$  in 1000 blocks
- Total cost 5500 I/Os.

Switch the inner and outer loop:

- Outer-loop: load chunk of 100 blocks of  $r$  to the main memory  
10 chunks, 100 blocks each
- Inner-loop: scan  $r$  in 500 blocks
- Total cost 6000 I/Os.

Conclusion: slight advantage in having the smaller relation on the outer loop.

## Examples: Cost Estimation, Merge Join

Consider relation  $r(A, B)$  with  $n_r = 1000$  blocks and relation  $s(A, C)$  with  $n_s = 500$  blocks. The memory had  $m = 101$  pages.

- The sorting cost of  $r$ : 4000 (two reads and two writes per block)
- The sorting cost of  $s$ : 2000 (two reads and two writes per block)
- Merge join: If  $A$  is the key of one relation, then the cost is 1500 blocks.
- Total cost 6500 I/Os.

Why linear time (when  $A$  is a key) merge join is not as good as quadratic time nested loop join in this case?

## Examples: Cost Estimation, Hash Join

Consider relation  $r(A, B)$  with  $n_r = 1000$  blocks and relation  $s(A, C)$  with  $n_s = 500$  blocks. The memory had  $m = 101$  pages. Let  $k = 100$ .

- The average size for each bucket is 10 blocks for relation  $r$  and 5 for  $s$ .
- Partition  $n$  and  $s$  (linear time):  $1500 \times 2 = 3000$  I/Os
- Since for each  $i \in [0, k)$ , the buckets of  $r$  and  $s$  can altogether fit in main memory, one-pass join: 1500 I/Os for loading.
- Total cost: 4500 I/Os.

# Overview



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE

1. Introduction
2. Physical-Query-Plan Operators
  - One-pass algorithms
  - Multi-pass algorithms
  - Index-based algorithms

# Index-based Algorithms

Application of Index in

- Selection
- Join

To avoid/reduce the number of table scan.

# Index-based Algorithms: Cost Estimation Assumptions

We first have to estimate sizes of the intermediate results.

- $V(A, r)$  number of distinct values of attribute  $A$  in relation  $r$ .
- $|r|$  number of tuples in relation  $r$ .
- $n_r$  number of blocks used to store relation  $r$ .


Important assumptions:

- Simple selection: All  $V(A, r)$  values are equally likely for attribute  $A$ ; in other words,  $|\sigma_{A=c}r| = |r|/V(A, r)$ .
- Selection involving inequality  $|\sigma_{A<c}r|$ : Common assumption that  $1/3$  will meet the condition.
- Complex conditions AND: use decompositions.
  - $|\sigma_{A=c \text{ and } B<d}r| = |r|/3V(A, r)$ .



# Index-based Algorithms: Selection

Sequential file:



The actual cost of  $\sigma_{A=v}r$  can be slightly larger than  $n_r/V(A, r)$ .

- The index is not kept entirely in main memory, some disk I/Os are needed to support the index lookup.
- Even when all the tuples with  $A = v$  might fit in  $b$  blocks, they could be spread over  $b + 1$  blocks because they don't start at the beginning of a block.
- The blocks can be not full, e.g., the  $B^+$  tree's leaf nodes can be not full.

Unordered file

- 33 ■ We assume that we need to visit  $|r|/V(A, r)$  blocks for answering  $\sigma_{A=v}r$ .

## Index-based Algorithms: Selection

Assume that for relation  $r(A, B)$ ,  $n_r = 1000$  and  $|r| = 20,000$ . Consider  $\sigma_{A=0}r$ . We ignore the cost accessing the index blocks in all cases.

- $r$  is sequential on  $A$  but we do not use the index 1000 I/Os.
- $r$  is unordered on  $A$  but we use an index on  $A$  20,000 I/Os.
- $r$  is sequential on  $A$ ,  $V(A, r) = 100$ , use an index 10 I/Os.
- $r$  is unordered on  $A$ ,  $V(A, r) = 10$ , use an index 2000 I/Os.
- $A$  is the key of  $r$ , use an index 1 I/O.

## Index-based Algorithms: Nested-Loop Join

$r(A, B) \bowtie s(A, C)$ :

- $r$  has  $|r|$  tuples on  $n_r$  blocks,  $s$  is stored in  $n_s$  blocks
- $s$  has an index on  $A$ ;

**for** each **tuple**  $t$  of  $r$

Get **all** tuples of  $s$  that can join **with**  $t$  using the index

**Complexity:** For each tuple  $t$  of  $r$ , an average of  $|s|/V(A, s)$  tuples should be retrieved, the cost is dependent on  $s$

- If  $s$  is sequential, for each  $t$  the cost is  $n_s/V(A, s)$  I/Os;
- Otherwise, the cost is  $|s|/V(A, s)$  I/Os.

## Index-based Algorithms: Nested-Loop Join


Example: Consider  $r(A, B)$  with  $n_r = 1000$  and  $s(A, C)$  with  $n_s = 500$  while ten tuples of either relation fit in one block. Therefore,  $|r| = 10,000$  and  $|s| = 5000$ . Assume that  $V(A, s) = 100$ . Suppose that  $s$  is ordered on  $A$  and there is a clustering index of  $s$  on  $A$ . Compute the cost of  $r \bowtie s$ .

- The number of I/Os in accessing  $s$  is  $10000 \times 500/100 = 50000$  I/Os.

If  $r$  is considerably smaller than  $s$  then the index-based nested-loop join would be better than nested-loop join.

# Index-based Algorithms: Merge Join

Both  $r$  and  $s$  are sequentially stored, each having a clustered index on  $A$ .



# Overview



**SCIENCE**  
DEPARTMENT OF  
COMPUTER SCIENCE

1. Introduction
2. Physical-Query-Plan Operators
  - One-pass algorithms
  - Multi-pass algorithms
  - Index-based algorithms
3. Query Optimization

Thank you for your attention!

Any questions?