



SOFTENG 351

Fundamentals of Database Systems

Basic SQL



Outline

- ▶ SQL Data Definition and Data Types
- ▶ Specifying Constraints in SQL
- ▶ Basic Retrieval Queries in SQL
- ▶ INSERT, DELETE, and UPDATE Statements in SQL
- ▶ Additional Features of SQL



Basic SQL

- ▶ SQL language
 - ▶ Considered one of the major reasons for the commercial success of relational databases.
- ▶ SQL
 - ▶ The origin of SQL is relational predicate calculus called tuple calculus, which was proposed initially as the language SQUARE.
 - ▶ Now popularly known as “Structured Query language”.
 - ▶ SQL is an informal or practical rendering of the **relational data model** with syntax.



SQL Standards

- ▶ SQL has gone through many standards: starting with SQL-86 (or SQL-1), SQL-92 is referred to as SQL-2.
- ▶ Later standards (from SQL-1999) are divided into **core** specification and specialized **extensions**. The extensions are implemented for different applications – such as data mining, data warehousing, multimedia etc.
- ▶ SQL-2006 added XML features; In 2008 added Object-oriented features.
- ▶ SQL-3 is the current standard which started with SQL-1999. It has not been fully implemented in any RDBMS yet.



SQL Data Definition and Data Types

- ▶ Terminology:
 - ▶ **Table, row, and column** used for relational model terms *relation, tuple, and attribute.*
- ▶ CREATE statement
 - ▶ Main SQL command for data definition
 - ▶ To create schemas, tables (relations), types, and domains,
 - ▶ and other constructs, e.g., views, assertions and triggers
- ▶ The language has features for : Data definition, Data Manipulation, Transaction control, Indexing, Security specification, Active databases, Multi-media, Distributed databases, etc.



Database Schema in SQL

- ▶ We cover the basic standard SQL syntax – there are variations in existing RDBMS systems
- ▶ **SQL schema**
 - ▶ Identified by a **schema name**
 - ▶ Includes an **authorization identifier** and **descriptors** for each element
- ▶ **Schema elements** include
 - ▶ Tables, types, constraints, views, domains, and other constructs (such as authorization grants)



Database Schema in SQL

- ▶ A database schema is created via the CREATE SCHEMA statement, e.g.,
- ▶ CREATE SCHEMA COMPANY AUTHORIZATION 'Jsmith';
- ▶ Each statement in SQL ends with a **semicolon**
- ▶ The privilege to create schemas, tables and other constructs must be explicitly granted to the relevant user accounts by the system administrator or DBA.



Catalog Concepts in SQL

▶ Catalog

- ▶ A named collection of schemas in an SQL environment
- ▶ A database installation typically have a default environment and schema.
- ▶ A catalog always contains a special schema called **INFORMATION_SCHEMA**, which provides information on all schemas in the catalog and all the element descriptors in these schemas.
- ▶ Schemas within the same catalog can share certain elements, such as type and domain definitions
- ▶ SQL also has the concept of a cluster of catalogs



The CREATE TABLE Command in SQL

- ▶ Specifying a new relation
 - ▶ Provide name of the table
 - ▶ Specify attributes, their types and initial constraints
- ▶ The attributes are specified:
 - ▶ Each attribute given a name, a data type to specify its domain of values
 - ▶ and possible attribute constraints, such as NOT NULL
- ▶ The key, entity integrity, and referential integrity constraints can be specified with the CREATE TABLE statement after the attributes are declared.
 - ▶ Or specified later using the ALTER TABLE statement



The CREATE TABLE Command in SQL

- ▶ Can attach the schema name to the relation name
 - ▶ CREATE TABLE COMPANY.EMPLOYEE ...
 - or
 - ▶ CREATE TABLE EMPLOYEE ...
- ▶ **Base tables (base relations)**
 - ▶ Relation and its tuples are actually created and stored as a file by the DBMS
- ▶ **Virtual relations (views)**
 - ▶ Created through the CREATE VIEW statement. May or may not correspond to an actual physical file.



COMPANY relational database schema

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------



One possible database state for the COMPANY relational database

EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

DEPT_LOCATIONS

Dnumber	Dlocation
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston



One possible database state for the COMPANY relational database

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

PROJECT

<u>Pname</u>	<u>Pnumber</u>	<u>Plocation</u>	<u>Dnum</u>
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	<u>Sex</u>	<u>Bdate</u>	<u>Relationship</u>
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse



SQL CREATE TABLE data definition statements for defining the COMPANY database

CREATE TABLE EMPLOYEE

(Fname	VARCHAR(15)	NOT NULL,
Minit	CHAR,	
Lname	VARCHAR(15)	NOT NULL,
Ssn	CHAR(9)	NOT NULL,
Bdate	DATE,	
Address	VARCHAR(30),	
Sex	CHAR,	
Salary	DECIMAL(10,2),	
Super_ssn	CHAR(9),	
Dno	INT	NOT NULL,

PRIMARY KEY (Ssn),

CREATE TABLE DEPARTMENT

(Dname	VARCHAR(15)	NOT NULL,
Dnumber	INT	NOT NULL,
Mgr_ssn	CHAR(9)	NOT NULL,
Mgr_start_date	DATE,	

PRIMARY KEY (Dnumber),

UNIQUE (Dname),

FOREIGN KEY (Mgr_ssn) **REFERENCES** EMPLOYEE(Ssn) ;

CREATE TABLE DEPT_LOCATIONS

(Dnumber	INT	NOT NULL,
Dlocation	VARCHAR(15)	NOT NULL,

PRIMARY KEY (Dnumber, Dlocation),

FOREIGN KEY (Dnumber) **REFERENCES** DEPARTMENT(Dnumber) ;

continued on next slide



SQL CREATE TABLE data definition statements for defining the COMPANY database

CREATE TABLE PROJECT

(Pname	VARCHAR(15)	NOT NULL,
Pnumber	INT	NOT NULL,
Plocation	VARCHAR(15),	
Dnum	INT	NOT NULL,

PRIMARY KEY (Pnumber),
UNIQUE (Pname),
FOREIGN KEY (Dnum) **REFERENCES** DEPARTMENT(Dnumber) ;

CREATE TABLE WORKS_ON

(Essn	CHAR(9)	NOT NULL,
Pno	INT	NOT NULL,
Hours	DECIMAL(3,1)	NOT NULL,

PRIMARY KEY (Essn, Pno),
FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn),
FOREIGN KEY (Pno) **REFERENCES** PROJECT(Pnumber) ;

CREATE TABLE DEPENDENT

(Essn	CHAR(9)	NOT NULL,
Dependent_name	VARCHAR(15)	NOT NULL,
Sex	CHAR,	
Bdate	DATE,	
Relationship	VARCHAR(8),	

PRIMARY KEY (Essn, Dependent_name),
FOREIGN KEY (Essn) **REFERENCES** EMPLOYEE(Ssn) ;



The CREATE TABLE Command in SQL

- ▶ In the previous example, some foreign keys might cause errors
 - ▶ Specified either via:
 - ▶ Circular references
 - ▶ Or because they refer to a table that has not yet been created
 - ▶ DBA's have ways to stop referential integrity enforcement to get around this problem.
 - ▶ These constraints can be left out of the initial CREATE TABLE statement, and then added later using the ALTER TABLE statement.



Attribute Data Types and Domains in SQL

- ▶ **Basic data types**
 - ▶ **Numeric data types**
 - ▶ **Integer numbers:** INTEGER, INT, and SMALLINT
 - ▶ **Floating-point (real) numbers:** FLOAT or REAL, and DOUBLE PRECISION
 - ▶ **Formatted numbers:** DECIMAL (i, j) or DEC (i, j), or NUMERIC (i, j)
 - i stands for *Precision*, the total number of digits in the value, i.e., on both sides of the decimal point
 - j stands for *Scale*, number of digits after the decimal point
 - ▶ **Boolean data type**
 - ▶ **Values of TRUE or FALSE (or NULL)**



Attribute Data Types and Domains in SQL

▶ **Character-string** data types

- ▶ Fixed length: CHAR (n), CHARACTER (n)
- ▶ Varying length : VARCHAR (n), CHAR VARYING (n), CHARACTER VARYING (n)

▶ **Bit-string** data types

- ▶ Fixed length: BIT (n)
- ▶ Varying length: BIT VARYING (n)

▶ **DATE** and **TIME** data types

- ▶ DATE has 10 positions, made of YEAR, MONTH, and DAY in the form YYYY-MM-DD
- ▶ TIME has at least 8 positions, made of HOUR, MINUTE, and SECOND in the form HH:MM:SS



Attribute Data Types and Domains in SQL

- ▶ Additional data types
 - ▶ **Timestamp** data type – Includes the DATE and TIME fields
 - ▶ Plus a minimum of six positions for decimal fractions of seconds
 - ▶ Optional WITH TIME ZONE qualifier
 - ▶ **INTERVAL** data type
 - ▶ Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp
 - ▶ **DATE, TIME, Timestamp, INTERVAL** data types can be **cast** or converted to string formats for comparison.



Attribute Data Types and Domains in SQL

► **Domain**

- ▶ Name used with the attribute specification
- ▶ Makes it easier to change the data type for a domain that is used by numerous attributes
- ▶ Improves schema readability
- ▶ Example:
 - ▶ `CREATE DOMAIN SSN_TYPE AS CHAR(9);`

► **TYPE**

- ▶ User Defined Types (UDTs) are supported for object-oriented applications. Uses the command:

```
CREATE TYPE
```



Specifying Constraints in SQL

Basic constraints:

- ▶ Relational Model has 3 basic constraint types that are supported in SQL:
 - ▶ Key constraint: A primary key value cannot be duplicated
 - ▶ Entity Integrity constraint: A primary key value cannot be null
 - ▶ Referential integrity constraints: The *foreign key* must have a value that is already present as a *primary key*, or may be null.



Specifying Attribute Constraints and Attribute Defaults

Basic constraints:

- ▶ Restrictions on attribute domains and NULL
- ▶ Default value of an attribute
 - ▶ **DEFAULT <value>**
- ▶ NOT NULL constraint
 - ▶ NULL is not permitted for a particular attribute
- ▶ For example,

```
CREATE TABLE EMPLOYEE
```

```
(...
```

```
Dno INT NOT NULL DEFAULT 1, ...)
```



Specifying Attribute Constraints and Attribute Defaults

Basic constraints:

- ▶ Constraints on individual tuples within a relation using the **CHECK** clause
 - ▶ `Dnumber INT NOT NULL CHECK (Dnumber > 0 AND Dnumber < 21);`
- ▶ The **CHECK** clause can also be used in conjunction with the **CREATE DOMAIN** statement.
- ▶ For example:

```
CREATE DOMAIN D_NUM AS INTEGER  
CHECK (D_NUM > 0 AND D_NUM < 21);
```



Specifying Key and Referential Integrity Constraints

▶ **PRIMARY KEY clause**

- ▶ Specifies one or more attributes that make up the **primary key of a relation**

▶ Dnumber INT PRIMARY KEY;

▶ **UNIQUE clause**

- ▶ Specifies alternate (secondary) keys (called **CANDIDATE keys** in the relational model).

▶ Dname VARCHAR (15) UNIQUE;



Specifying Key and Referential Integrity Constraints

- ▶ **FOREIGN KEY** clause, e.g.,

```
FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
    ON DELETE SET NULL          ON UPDATE CASCADE,
```

- ▶ Default action for violation is to reject update operation
- ▶ Attach **referential triggered action** clause
 - ▶ Options include SET NULL, CASCADE, and SET DEFAULT
 - ▶ Action taken by the DBMS for SET NULL or SET DEFAULT is the same for both ON DELETE and ON UPDATE
 - ▶ CASCADE option suitable for “relationship” relations



Giving Names to Constraints

- ▶ Using the Keyword **CONSTRAINT**
 - ▶ Name a constraint
 - ▶ Useful for later altering
- ▶ For example:

CONSTRAINT EMPPK

PRIMARY KEY (Ssn),

CONSTRAINT EMPSUPERFK

FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)

ON DELETE SET NULL

ON UPDATE CASCADE,

CONSTRAINT EMPDEPTFK

FOREIGN KEY(Dno) REFERENCES DEPARTMENT(Dnumber)

ON DELETE SET DEFAULT

ON UPDATE CASCADE);



Default attribute values and referential integrity triggered action specification

```
CREATE TABLE EMPLOYEE
(
    ...,
    Dno      INT      NOT NULL      DEFAULT 1,
    CONSTRAINT EMPPK
        PRIMARY KEY (Ssn),
    CONSTRAINT EMPSUPERFK
        FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
            ON DELETE SET NULL      ON UPDATE CASCADE,
    CONSTRAINT EMPDEPTFK
        FOREIGN KEY(Dno) REFERENCES DEPARTMENT(Dnumber)
            ON DELETE SET DEFAULT  ON UPDATE CASCADE);
CREATE TABLE DEPARTMENT
(
    ...,
    Mgr_ssn CHAR(9)      NOT NULL      DEFAULT '888665555',
    ...,
    CONSTRAINT DEPTPK
        PRIMARY KEY(Dnumber),
    CONSTRAINT DEPTSK
        UNIQUE (Dname),
    CONSTRAINT DEPTMGRFK
        FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
            ON DELETE SET DEFAULT  ON UPDATE CASCADE);
CREATE TABLE DEPT_LOCATIONS
(
    ...,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
        ON DELETE CASCADE      ON UPDATE CASCADE);
```



Specifying Constraints on Tuples Using CHECK

- ▶ Additional Constraints on individual tuples within a relation are also possible using **CHECK**

- ▶ **CHECK clauses at the end of a CREATE TABLE statement**

- ▶ **Apply to each tuple individually**

```
CHECK (Dept_create_date <=
        Mgr_start_date);
```

- ▶ These are **row-based constraints**, because they apply to each row individually and are checked whenever a row is inserted or modified.



Basic Retrieval Queries in SQL

- ▶ **SELECT statement**
 - ▶ One basic statement for retrieving information from a database
- ▶ Important distinction between the practical **SQL** model and the formal **Relational** model discussed in topic 4.
 - ▶ SQL allows a table to have two or more tuples that are identical in all their attribute values
 - ▶ Unlike relational model (relational model is strictly set-theory based)
 - ▶ SQL table is not a set of tuples, but a Multiset (or a bag)
 - ▶ Some SQL relations are constrained to be sets because a key constraint has been declared.



The SELECT-FROM-WHERE Structure of Basic SQL Queries

► Basic form of the SELECT statement:

```
SELECT      <attribute list>
FROM        <table list>
WHERE       <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.



The SELECT-FROM-WHERE Structure of Basic SQL Queries

- ▶ Basic logical comparison operators for comparing attribute values
 - ▶ $=$, $<$, \leq , $>$, \geq , and \neq
- ▶ **Projection attributes**
 - ▶ Attributes whose values are to be retrieved
- ▶ **Selection condition**
 - ▶ Boolean condition that must be true for any retrieved tuple. Selection conditions include join conditions when multiple relations are involved.
- ▶ Consider an implicit tuple variable (or iterator) looping over individual tuples and evaluating against the condition.



Basic Retrieval Queries

<u>Bdate</u>	<u>Address</u>
1965-01-09	731 Fondren, Houston, TX

<u>Fname</u>	<u>Lname</u>	<u>Address</u>
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

Query 0. Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0: **SELECT** Bdate, Address
 FROM EMPLOYEE
 WHERE Fname='John' **AND** Minit='B' **AND** Lname='Smith';

Query 1. Retrieve the name and address of all employees who work for the 'Research' department.

Q1: **SELECT** Fname, Lname, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE Dname='Research' **AND** Dnumber=Dno;



Basic Retrieval Queries

(c)

<u>Pnumber</u>	<u>Dnum</u>	<u>Lname</u>	<u>Address</u>	<u>Bdate</u>
10	4	Wallace	291Berry, Bellaire, TX	1941-06-20
30	4	Wallace	291Berry, Bellaire, TX	1941-06-20

Query 2. For every project located in ‘Stafford’, list the project number, the controlling department number, and the department manager’s last name, address, and birth date.

Q2: **SELECT** Pnumber, Dnum, Lname, Address, Bdate
 FROM PROJECT, DEPARTMENT, EMPLOYEE
 WHERE Dnum=Dnumber **AND** Mgr_ssn=Ssn **AND**
 Plocation=‘Stafford’;



Ambiguous Attribute Names

- ▶ Same name can be used for two (or more) attributes in different relations
 - ▶ As long as the attributes are in different relations
 - ▶ Must **qualify** the attribute name with the relation name to prevent ambiguity

Q1A: **SELECT** Fname, EMPLOYEE.Name, Address
 FROM EMPLOYEE, DEPARTMENT
 WHERE DEPARTMENT.Name='Research' **AND**
 DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;



Aliasing, and Renaming

▶ Aliases or tuple variables

- ▶ Declare alternative relation names E and S to refer to the EMPLOYEE relation twice in a query:

Query 8. For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

- ▶

```
SELECT E.Fname, E.Lname, S.Fname, S.Lname
FROM EMPLOYEE AS E, EMPLOYEE AS S
WHERE E.Super_ssn=S.Ssn;
```

 - ▶ Recommended practice to abbreviate names and to prefix same or similar attribute from multiple tables.



Aliasing, Renaming and Tuple Variables

- ▶ The attribute names can also be renamed, e.g.,

EMPLOYEE AS E(Fn, Mi, Ln, Ssn, Bd,
Addr, Sex, Sal, Sssn, Dno)

- ▶ Note that the relation EMPLOYEE now has a variable name E which corresponds to a tuple variable
- ▶ WE can use alias-naming or renaming mechanism in any SQL query to specify tuple variables. E.g.,

Q1B: **SELECT** E.Fname, E.LName, E.Address
 FROM EMPLOYEE **AS** E, DEPARTMENT **AS** D
 WHERE D.DName = 'Research' **AND** D.Dnumber = E.Dno;



Unspecified WHERE Clause and Use of the Asterisk

- ▶ Missing WHERE clause
 - ▶ Indicates no condition on tuple selection
- ▶ Effect is a CROSS PRODUCT
 - ▶ Result is all possible tuple combinations result

Queries 9 and 10. Select all EMPLOYEE Ssns (Q9) and all combinations of EMPLOYEE Ssn and DEPARTMENT Dname (Q10) in the database.

Q9: **SELECT** Ssn
 FROM EMPLOYEE;

Q10: **SELECT** Ssn, Dname
 FROM EMPLOYEE, DEPARTMENT;



Unspecified WHERE Clause and Use of the Asterisk

- ▶ Specify an asterisk (*)
 - ▶ Retrieve all the attribute values of the selected tuples
 - ▶ The * can be prefixed by the relation name, e.g., EMPLOYEE.* refers to all attributes of EMPLOYEE table

Q1C: **SELECT** *
FROM EMPLOYEE
WHERE Dno=5;

Q1D: **SELECT** *
FROM EMPLOYEE, DEPARTMENT
WHERE Dname='Research' **AND** Dno=Dnumber;

Q10A: **SELECT** *
FROM EMPLOYEE, DEPARTMENT;



Tables as Sets in SQL

- ▶ SQL does not automatically eliminate duplicate tuples in query results
- ▶ For aggregate operations duplicates must be accounted for
- ▶ Use the keyword **DISTINCT** in the SELECT clause
 - ▶ Only distinct tuples should remain in the result

Query 11. Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

Q11: **SELECT** **ALL** **Salary**
 FROM **EMPLOYEE;**

Q11A: **SELECT** **DISTINCT** **Salary**
 FROM **EMPLOYEE;**



Tables as Sets in SQL

- ▶ SQL has directly incorporated some of set operations
 - ▶ UNION, EXCEPT (i.e., set difference), INTERSECT

Query 4. Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A: (SELECT DISTINCT Pnumber
      FROM PROJECT, DEPARTMENT, EMPLOYEE
      WHERE Dnum=Dnumber AND Mgr_ssn=Ssn
            AND Lname='Smith')
      UNION
      (SELECT DISTINCT Pnumber
      FROM PROJECT, WORKS_ON, EMPLOYEE
      WHERE Pnumber=Pno AND Essn=Ssn
            AND Lname='Smith');
```

- ▶ Corresponding multiset operations: UNION ALL, EXCEPT ALL, INTERSECT ALL)
- ▶ Type compatibility is needed for these operations to be valid



Substring Pattern Matching and Arithmetic Operators

- ▶ **LIKE** comparison operator
 - ▶ Used for string pattern matching
 - ▶ % replaces an arbitrary number of zero or more characters
 - ▶ underscore (_) replaces a single character
 - ▶ Examples: **WHERE** Address **LIKE** '%Houston,TX%';
 - ▶ **WHERE** Ssn **LIKE** '_ _ I _ _ 8901';
- ▶ **BETWEEN** comparison operator, e.g.:
WHERE(Salary **BETWEEN** 30000 **AND** 40000)
AND Dno = 5;



Arithmetic Operations

- ▶ Standard arithmetic operators:
 - ▶ Addition (+), subtraction (-), multiplication (*), and division (/) may be included as a part of **SELECT**
- ▶ **Query 13.** Show the resulting salaries if every employee working on the ‘ProductX’ project is given a 10 percent raise.

```
SELECT E.Fname, E.Lname, I.I * E.Salary AS Increased_sal  
FROM EMPLOYEE AS E,WORKS_ON AS W,PROJECT AS P  
WHERE E.Ssn=W.Essn AND W.Pno=P.Pnumber AND  
P.Pname='ProductX';
```



Ordering of Query Results

▶ Use **ORDER BY** clause

```
Q15:   SELECT      D.Dname, E.Lname, E.Fname, P.Pname  
          FROM        DEPARTMENT AS D, EMPLOYEE AS E, WORKS_ON AS W,  
                      PROJECT AS P  
         WHERE       D.Dnumber = E.Dno AND E.Ssn = W.Essn AND W.Pno =  
                      P.Pnumber  
         ORDER BY    D.Dname, E.Lname, E.Fname;
```

- ▶ The default order is ascending order of values
- ▶ Keyword **DESC** to set a descending order of values
- ▶ Keyword **ASC** to specify ascending order explicitly
- ▶ Typically placed at the end of the query

```
ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC
```



Basic SQL Retrieval Query Block

- ▶ A simple retrieval query in SQL can consists of up to four clauses,
 - ▶ only the first two SELECT and FROM are mandatory.

```
SELECT      <attribute list>
FROM        <table list>
[ WHERE     <condition> ]
[ ORDER BY   <attribute list> ];
```

- ▶ SELECT – lists the attributes to be retrieved
- ▶ FROM – specifies all relations (tables) needed
- ▶ WHERE – identifies conditions for the selection
- ▶ ORDERBY – specifies an order for displaying the results



INSERT, DELETE, and UPDATE Statements in SQL

- ▶ Three commands used to modify the database, i.e.,
 - ▶ INSERT, DELETE, and UPDATE
- ▶ **INSERT** typically inserts a single tuple (row) into a relation (table)
- ▶ **UPDATE** modifies attribute values of a number of tuples (rows) in a relation (table) that satisfy the condition
- ▶ **DELETE** removes a number of tuples (rows) in a relation (table) that satisfy the condition
- ▶ Both UPDATE And DELETE may propagate to tuples in other relations if referential triggered actions are specified in the referential integrity constraints of the DDL.



INSERT

- ▶ In its simplest form, it is used to add one or more tuples to a relation
- ▶ Attribute values should be listed in the same order as the attributes were specified in the **CREATE TABLE** command
- ▶ Constraints on data types are observed automatically
- ▶ Any integrity constraints as a part of the DDL specification are enforced



The INSERT Command

- ▶ Specify the relation name and a list of values for the tuple. All values including nulls are supplied.

```
U1:  INSERT INTO EMPLOYEE  
      VALUES ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98  
              Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

- ▶ The variation below inserts multiple tuples where a new table is loaded values from the result of a query.

```
U3B:  INSERT INTO WORKS_ON_INFO ( Emp_name, Proj_name,
                                  Hours_per_week )
        SELECT E.Lname, P.Pname, W.Hours
        FROM PROJECT P, WORKS_ON W, EMPLOYEE E
       WHERE P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```



Bulk Loading of TABLEs

- ▶ Another variation of **INSERT** is used for bulk-loading of several tuples into tables
- ▶ A new table TNEW can be created with the same attributes as T and using LIKE and DATA in the syntax, it can be loaded with entire data.
- ▶ For example:

```
CREATE TABLE D5EMPS LIKE EMPLOYEE  
(SELECT E.*  
      FROM    EMPLOYEE AS E  
      WHERE   E.Dno=5)  
WITH DATA;
```



DELETE

- ▶ Removes tuples from a relation
 - ▶ Includes a WHERE-clause to select the tuples to be deleted
 - ▶ Referential integrity should be enforced
 - ▶ Tuples are deleted from only one *table* at a time (unless CASCADE is specified on a referential integrity constraint)
 - ▶ A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table
 - ▶ The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause



The DELETE Command

- ▶ Removes tuples from a relation
 - ▶ Includes a WHERE clause to select the tuples to be deleted. The number of tuples deleted will vary.

U4A:	DELETE FROM	EMPLOYEE
	WHERE	Lname='Brown';
U4B:	DELETE FROM	EMPLOYEE
	WHERE	Ssn='123456789';
U4C:	DELETE FROM	EMPLOYEE
	WHERE	Dno=5;
U4D:	DELETE FROM	EMPLOYEE;



UPDATE

- ▶ Used to modify attribute values of one or more selected tuples
- ▶ A WHERE-clause selects the tuples to be modified
- ▶ An additional SET-clause specifies the attributes to be modified and their new values
- ▶ Each command modifies tuples *in the same relation*
- ▶ Referential integrity specified as part of DDL specification is enforced



The UPDATE Command

- ▶ An additional SET-clause specifies the attributes to be modified and their new values
- ▶ For example:
- ▶ Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively

U5: **UPDATE** **PROJECT**
 SET **Plocation = 'Bellaire', Dnum = 5**
 WHERE **Pnumber = 10;**



The UPDATE Command

- ▶ For example: Give all employees in the 'Research' department a 10% raise in salary.

```
U6: UPDATE EMPLOYEE  
      SET SALARY = SALARY * 1.1  
      WHERE DNO IN (SELECT DNUMBER  
                     FROM DEPARTMENT  
                     WHERE DNAME='Research')
```

- ▶ In this request, the modified SALARY value depends on the original SALARY value in each tuple
 - ▶ The reference to the SALARY attribute on the right of = refers to the old SALARY value before modification
 - ▶ The reference to the SALARY attribute on the left of = refers to the new SALARY value after modification



Additional Features of SQL

- ▶ Techniques for specifying complex retrieval queries
- ▶ Writing programs in various programming languages that include SQL statements: Embedded and dynamic SQL, SQL/CLI (Call Level Interface) and its predecessor ODBC, SQL/PSM (Persistent Stored Module)
- ▶ Set of commands for specifying physical database design parameters, file structures for relations, and access paths, e.g., CREATE INDEX



Additional Features of SQL

- ▶ Transaction control commands
- ▶ Specifying the granting and revoking of privileges to users
- ▶ Constructs for creating triggers
- ▶ Enhanced relational systems known as object-relational define relations as classes. Abstract data types (called User Defined Types- UDTs) are supported with CREATE TYPE
- ▶ New technologies such as XML and OLAP are added to versions of SQL



Summary

- ▶ **SQL**
 - ▶ A Comprehensive language for relational database management
 - ▶ Data definition, queries, updates, constraint specification, and view definition
- ▶ **Covered :**
 - ▶ Data definition commands for creating tables
 - ▶ Commands for constraint specification
 - ▶ Simple retrieval queries
 - ▶ Database update commands